

## Relational Algebra, SQL, Triggers, and Active Database

1

## Relational Query Languages

- Languages for describing queries on a relational database
- *Structured Query Language* (SQL)
  - Predominant application-level query language
  - Declarative
- *Relational Algebra*
  - Intermediate language used within DBMS
  - Procedural

2

## What is an Algebra?

- A language based on operators and a domain of values
- Operators map values taken from the domain into other domain values
- Hence, an expression involving operators and arguments produces a value in the domain
- When the domain is a set of all relations (and the operators are as described later), we get the *relational algebra*
- We refer to the expression as a *query* and the value produced as the *query result*

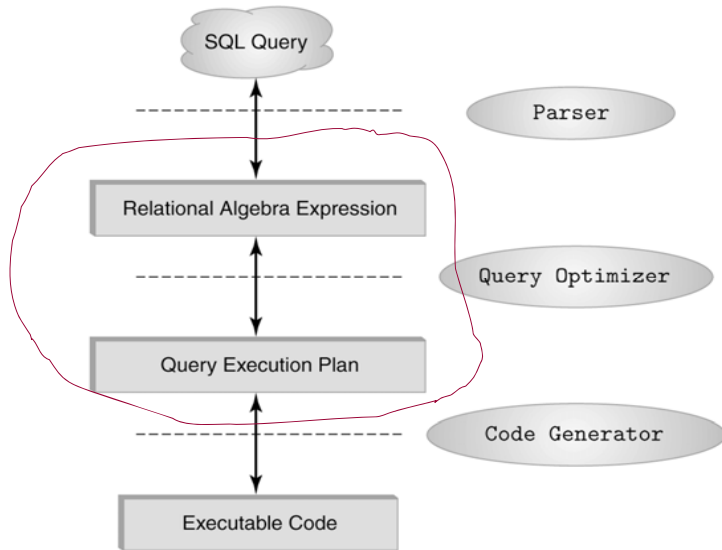
3

## Relational Algebra

- *Domain*: set of relations
- *Basic operators*: select, project, union, set difference, Cartesian product
- *Derived operators*: set intersection, division, join
- *Procedural*: Relational expression specifies query by describing an algorithm (the sequence in which operators are applied) for determining the result of an expression

4

## The Role of Relational Algebra in a DBMS



5

## Schema for Student Registration System

Student (*Id*, *Name*, *Addr*, *Status*)

Professor (*Id*, *Name*, *DeptId*)

Course (*DeptId*, *CrsCode*, *CrsName*, *Descr*)

Transcript (*StudId*, *CrsCode*, *Semester*, *Grade*)

Teaching (*ProfId*, *CrsCode*, *Semester*)

Department (*DeptId*, *Name*)

6

## Select Operator

- Produce table containing subset of rows of argument table satisfying condition

$\sigma_{condition}(relation)$

- Example:

Person

<i>Id</i>	<i>Name</i>	<i>Address</i>	<i>Hobby</i>
1123	John	123 Main	stamps
1123	John	123 Main	coins
5556	Mary	7 Lake Dr	hiking
9876	Bart	5 Pine St	stamps

$\sigma_{Hobby='stamps'}(Person)$

<i>Id</i>	<i>Name</i>	<i>Address</i>	<i>Hobby</i>
1123	John	123 Main	stamps
9876	Bart	5 Pine St	stamps

7

## Selection Condition

- Operators:  $<$ ,  $\leq$ ,  $\geq$ ,  $>$ ,  $=$ ,  $\neq$
- Simple selection condition:
  - $\langle attribute \rangle operator \langle constant \rangle$
  - $\langle attribute \rangle operator \langle attribute \rangle$
- $\langle condition \rangle AND \langle condition \rangle$
- $\langle condition \rangle OR \langle condition \rangle$
- NOT  $\langle condition \rangle$

8

## Selection Condition - Examples

- $\sigma_{Id > 3000 \text{ OR } Hobby = 'hiking'}(Person)$
- $\sigma_{Id > 3000 \text{ AND } Id < 3999}(Person)$
- $\sigma_{NOT(Hobby = 'hiking')}(Person)$
- $\sigma_{Hobby \neq 'hiking'}(Person)$

9

## Project Operator

- Produces table containing subset of columns of argument table

$$\pi_{attribute\ list}(relation)$$

- Example:

Person

<i>Id</i>	<i>Name</i>	<i>Address</i>	<i>Hobby</i>
1123	John	123 Main	stamps
1123	John	123 Main	coins
5556	Mary	7 Lake Dr	hiking
9876	Bart	5 Pine St	stamps

$\pi_{Name, Hobby}(Person)$

<i>Name</i>	<i>Hobby</i>
John	stamps
John	coins
Mary	hiking
Bart	stamps

10

## Project Operator

- Example:

Person

<i>Id</i>	<i>Name</i>	<i>Address</i>	<i>Hobby</i>
1123	John	123 Main	stamps
1123	John	123 Main	coins
5556	Mary	7 Lake Dr	hiking
9876	Bart	5 Pine St	stamps

$\pi_{Name, Address}(Person)$

<i>Name</i>	<i>Address</i>
John	123 Main
Mary	7 Lake Dr
Bart	5 Pine St

Result is a table (no duplicates); can have fewer tuples than the original

11

## Expressions

$$\pi_{Id, Name}(\sigma_{Hobby = 'stamps' \text{ OR } Hobby = 'coins'}(Person))$$

<i>Id</i>	<i>Name</i>	<i>Address</i>	<i>Hobby</i>
1123	John	123 Main	stamps
1123	John	123 Main	coins
5556	Mary	7 Lake Dr	hiking
9876	Bart	5 Pine St	stamps

Person

<i>Id</i>	<i>Name</i>
1123	John
9876	Bart

Result

12

## Set Operators

- Relation is a set of tuples, so set operations should apply:  $\cap$ ,  $\cup$ ,  $-$  (set difference)
- Result of combining two relations with a set operator is a relation  $\Rightarrow$  all its elements must be tuples having same structure
- Hence, scope of set operations limited to *union compatible relations*

13

## Union Compatible Relations

- Two relations are *union compatible* if
  - Both have same number of columns
  - Names of attributes are the same in both
  - Attributes with the same name in both relations have the same domain
- Union compatible relations can be combined using *union*, *intersection*, and *set difference*

14

## Example

Tables:

Person (*SSN, Name, Address, Hobby*)

Professor (*Id, Name, Office, Phone*)

are not union compatible.

But

$\pi_{Name}(\text{Person})$  and  $\pi_{Name}(\text{Professor})$

are union compatible so

$\pi_{Name}(\text{Person}) - \pi_{Name}(\text{Professor})$

makes sense.

15

## Cartesian Product

- If  $R$  and  $S$  are two relations,  $R \times S$  is the set of all concatenated tuples  $\langle x, y \rangle$ , where  $x$  is a tuple in  $R$  and  $y$  is a tuple in  $S$ 
  - $R$  and  $S$  need not be union compatible
- $R \times S$  is expensive to compute:
  - Factor of two in the size of each row
  - Quadratic in the number of rows

<i>A</i>	<i>B</i>	<i>C</i>	<i>D</i>
x1	x2	y1	y2
x3	x4	y3	y4

$R$                    $S$

<i>A</i>	<i>B</i>	<i>C</i>	<i>D</i>
x1	x2	y1	y2
x1	x2	y3	y4
x3	x4	y1	y2
x3	x4	y3	y4

$R \times S$

16

## Renaming

- Result of expression evaluation is a relation
- Attributes of relation must have distinct names. This is not guaranteed with Cartesian product
  - e.g., suppose in previous example  $a$  and  $c$  have the same name
- Renaming operator tidies this up. To assign the names  $A_1, A_2, \dots, A_n$  to the attributes of the  $n$  column relation produced by expression  $expr$  use  $expr [A_1, A_2, \dots, A_n]$

17

## Example

Transcript ( $StudId$ ,  $CrsCode$ ,  $Semester$ ,  $Grade$ )

Teaching ( $ProfId$ ,  $CrsCode$ ,  $Semester$ )

$\pi_{StudId, CrsCode} (Transcript) [StudId, CrsCode1]$   
 $\times \pi_{ProfId, CrsCode} (Teaching) [ProfId, CrsCode2]$

This is a relation with 4 attributes:

$StudId, CrsCode1, ProfId, CrsCode2$

18

## Derived Operation: Join

A (general or theta) join of  $R$  and  $S$  is the expression

$$R \bowtie_{join-condition} S$$

where  $join-condition$  is a conjunction of terms:

$$A_i \text{ oper } B_j$$

in which  $A_i$  is an attribute of  $R$ ;  $B_j$  is an attribute of  $S$ ; and  $oper$  is one of  $=, <, >, \geq, \neq, \leq$ .

The meaning is:

$$\sigma_{join-condition'} (R \times S)$$

where  $join-condition$  and  $join-condition'$  are the same, except for possible renamings of attributes (next)

19

## Join and Renaming

- **Problem:**  $R$  and  $S$  might have attributes with the same name – in which case the Cartesian product is not defined
- **Solutions:**
  1. Rename attributes prior to forming the product and use new names in  $join-condition'$ .
  2. Qualify common attribute names with relation names (thereby disambiguating the names). For instance: Transcript. $CrsCode$  or Teaching. $CrsCode$ 
    - This solution is nice, but doesn't always work: consider

$$R \bowtie_{join\_condition} R$$

In  $R.A$ , how do we know which  $R$  is meant?

20

## Theta Join – Example

Employee(Name,Id,MngrId,Salary)

Manager(Name,Id,Salary)

Output the names of all employees that earn more than their managers.

$\pi_{\text{Employee.Name}} (\text{Employee} \bowtie_{\text{MngrId=Id AND Salary>Salary}} \text{Manager})$

The join yields a table with attributes:

Employee.Name, Employee.Id, Employee.Salary, MngrId  
Manager.Name, Manager.Id, Manager.Salary

21

## Equijoin Join - Example

*Equijoin*: Join condition is a conjunction of *equalities*.

$\pi_{\text{Name,CrsCode}} (\text{Student} \bowtie_{\text{Id=StudId}} \sigma_{\text{Grade='A'}} (\text{Transcript}))$

Student				Transcript			
Id	Name	Addr	Status	StudId	CrsCode	Sem	Grade
111	John	.....	.....	111	CSE305	S00	B
222	Mary	.....	.....	222	CSE306	S99	A
333	Bill	.....	.....	333	CSE304	F99	A
444	Joe	.....	.....				

Mary	CSE306
Bill	CSE304

*The equijoin is used very frequently since it combines related data in different relations.*

22

## Natural Join

- Special case of equijoin:
  - join condition equates *all* and *only* those attributes with the same name (condition doesn't have to be explicitly stated)
  - duplicate columns eliminated from the result

Transcript (StudId, CrsCode, Sem, Grade)
Teaching (ProfId, CrsCode, Sem)

Transcript  $\bowtie$  Teaching =

$\pi_{\text{StudId, Transcript.CrsCode, Transcript.Sem, Grade, ProfId}}$   
 $(\text{Transcript} \bowtie_{\text{CrsCode=CrsCode AND Sem=Sem}} \text{Teaching})$   
 $[\text{StudId, CrsCode, Sem, Grade, ProfId}]$

23

## Natural Join (cont'd)

- More generally:

$$R \bowtie S = \pi_{\text{attr-list}} (\sigma_{\text{join-cond}} (R \times S))$$

where

$\text{attr-list} = \text{attributes}(R) \cup \text{attributes}(S)$   
 (duplicates are eliminated) and  $\text{join-cond}$  has the form:

$$A_1 = A_1 \text{ AND } \dots \text{ AND } A_n = A_n$$

where

$$\{A_1 \dots A_n\} = \text{attributes}(R) \cap \text{attributes}(S)$$

24

## Natural Join Example

- List all Ids of students who took at least two different courses:

$\pi_{StudId} ( \sigma_{CrsCode \neq CrsCode2} ($   
 Transcript  $\bowtie$   
 Transcript [StudId, CrsCode2, Sem2, Grade2] ))

We don't want to join on *CrsCode*, *Sem*, and *Grade* attributes, hence renaming!


25

## Division

- Goal: Produce the tuples in one relation,  $r$ , that match *all* tuples in another relation,  $s$ 
  - $r (A_1, \dots A_n, B_1, \dots B_m)$
  - $s (B_1 \dots B_m)$
  - $r/s$ , with attributes  $A_1, \dots A_n$ , is the set of all tuples  $\langle a \rangle$  such that for every tuple  $\langle b \rangle$  in  $s$ ,  $\langle a, b \rangle$  is in  $r$
- Can be expressed in terms of projection, set difference, and cross-product

26

## Division (cont'd)



27

## Division - Example

- List the Ids of students who have passed all courses that were taught in spring 2000
- Numerator:**
  - $StudId$  and  $CrsCode$  for every course passed by every student:
$$\pi_{StudId, CrsCode} ( \sigma_{Grade \neq 'F'} (Transcript) )$$
- Denominator:**
  - $CrsCode$  of all courses taught in spring 2000
$$\pi_{CrsCode} ( \sigma_{Semester='S2000'} (Teaching) )$$
- Result is *numerator/denominator*

28

## Schema for Student Registration System

Student (*Id, Name, Addr, Status*)

Professor (*Id, Name, DeptId*)

Course (*DeptId, CrsCode, CrsName, Descr*)

Transcript (*StudId, CrsCode, Semester, Grade*)

Teaching (*ProfId, CrsCode, Semester*)

Department (*DeptId, Name*)

29

## Query Sublanguage of SQL

```
SELECT C.CrsName
FROM Course C
WHERE C.DeptId = 'CS'
```

- *Tuple variable* C ranges over rows of Course.
- Evaluation strategy:
  - FROM clause produces Cartesian product of listed tables
  - WHERE clause assigns rows to C in sequence and produces table containing only rows satisfying condition
  - SELECT clause retains listed columns
- Equivalent to:  $\pi_{CrsName} \sigma_{DeptId='CS'}(Course)$

30

## Join Queries

```
SELECT C.CrsName
FROM Course C, Teaching T
WHERE C.CrsCode=T.CrsCode AND T.Semester='S2000'
```

- List CS courses taught in S2000
- Tuple variables clarify meaning.
- Join condition “C.CrsCode=T.CrsCode”
  - relates facts to each other
- Selection condition “T.Semester='S2000'”
  - eliminates irrelevant rows
- Equivalent (using natural join) to:

$\pi_{CrsName}(Course \bowtie \sigma_{Semester='S2000'}(Teaching))$

$\pi_{CrsName}(\sigma_{Sem='S2000'}(Course \bowtie Teaching))$

31

## Correspondence Between SQL and Relational Algebra

```
SELECT C.CrsName
FROM Course C, Teaching T
WHERE C.CrsCode = T.CrsCode AND T.Semester = 'S2000'
```

Also equivalent to:

$\pi_{CrsName} \sigma_{C\_CrsCode=T\_CrsCode \text{ AND } Semester='S2000'}$   
 $(Course [C\_CrsCode, DeptId, CrsName, Desc]$   
 $\times Teaching [ProfId, T\_CrsCode, Semester])$

- This is the simplest evaluation algorithm for SELECT.
- Relational algebra expressions are procedural.
  - Which of the two equivalent expressions is more easily evaluated?

32



## Self-join Queries

Find Ids of all professors who taught at least two courses in the same semester:

```
SELECT T1.ProfId
FROM Teaching T1, Teaching T2
WHERE T1.ProfId = T2.ProfId
      AND T1.Semester = T2.Semester
      AND T1.CrsCode <> T2.CrsCode
```

*Tuple variables are essential in this query!*

Equivalent to:

```
 $\pi_{ProfId}(\sigma_{T1.CrsCode \neq T2.CrsCode}(\text{Teaching}[ProfId, T1.CrsCode, Semester] \bowtie \text{Teaching}[ProfId, T2.CrsCode, Semester]))$ 
```

33

## Duplicates

- Duplicate rows not allowed in a relation
- However, duplicate elimination from query result is costly and not done by default; must be explicitly requested:

```
SELECT DISTINCT .....
FROM .....
```

34

## Use of Expressions

Equality and comparison operators apply to strings (based on lexical ordering)

```
WHERE S.Name < 'P'
```

Concatenate operator applies to strings

```
WHERE S.Name || '--' || S.Address = .....
```

Expressions can also be used in SELECT clause:

```
SELECT S.Name || '--' || S.Address AS NmAdd
FROM Student S
```

35

## Set Operators

- SQL provides UNION, EXCEPT (set difference), and INTERSECT for union compatible tables
- Example: Find all professors in the CS Department and all professors that have taught CS courses

```
(SELECT P.Name
FROM Professor P, Teaching T
WHERE P.Id=T.ProfId AND T.CrsCode LIKE 'CS%')
UNION
(SELECT P.Name
FROM Professor P
WHERE P.DeptId = 'CS')
```

36

## Nested Queries

List all courses that were not taught in S2000

```
SELECT C.CrsName
FROM Course C
WHERE C.CrsCode NOT IN
  (SELECT T.CrsCode  --subquery
   FROM Teaching T
   WHERE T.Sem = 'S2000')
```

Evaluation strategy: subquery evaluated once to produces set of courses taught in S2000. Each row (as C) tested against this set.

37

## Correlated Nested Queries

Output a row  $\langle prof, dept \rangle$  if  $prof$  has taught a course in  $dept$ .

```
SELECT P.Name, D.Name      --outer query
FROM Professor P, Department D
WHERE P.Id IN
  -- set of all ProfId's who have taught a course in D.DeptId
  (SELECT T.ProfId          --subquery
   FROM Teaching T, Course C
   WHERE T.CrsCode=C.CrsCode AND
         C.DeptId=D.DeptId  --correlation
  )
```

38


## Correlated Nested Queries (con't)

- Tuple variables T and C are *local* to subquery
- Tuple variables P and D are *global* to subquery
- *Correlation*: subquery uses a global variable, D
- The value of D.DeptId parameterizes an evaluation of the subquery
- Subquery must (at least) be re-evaluated for each distinct value of D.DeptId
- *Correlated queries can be expensive to evaluate*

39

## Division in SQL

- *Query type*: Find the subset of items in one set that are related to *all* items in another set
- *Example*: Find professors who taught courses in *all* departments
  - Why does this involve division?



$$\pi_{\text{ProfId}, \text{DeptId}}(\text{Teaching} \bowtie \text{Course}) / \pi_{\text{DeptId}}(\text{Department})$$

40

## Division in SQL

- *Strategy for implementing division in SQL:*
  - Find set, A, of all departments in which a particular professor,  $p$ , has taught a course
  - Find set, B, of all departments
  - Output  $p$  if  $A \supseteq B$ , or, equivalently, if  $B-A$  is empty

41

## Division – SQL Solution

```
SELECT P.Id
FROM Professor P
WHERE NOT EXISTS
    (SELECT D.DeptId      -- set B of all dept Ids
      FROM Department D
      EXCEPT
      SELECT C.DeptId     -- set A of dept Ids of depts in
                           -- which P taught a course
      FROM Teaching T, Course C
      WHERE T.ProfId=P.Id -- global variable
            AND T.CrsCode=C.CrsCode)
```

42

## Aggregates

- Functions that operate on sets:
  - COUNT, SUM, AVG, MAX, MIN
- Produce numbers (not tables)
- Not part of relational algebra (but not hard to add)

```
SELECT COUNT(*)      SELECT MAX (Salary)
FROM Professor P      FROM Employee E
```

43

## Aggregates (cont'd)

Count the number of courses taught in S2000

```
SELECT COUNT (T.CrsCode)
FROM Teaching T
WHERE T.Semester = 'S2000'
```

But if multiple sections of same course are taught, use:

```
SELECT COUNT (DISTINCT T.CrsCode)
FROM Teaching T
WHERE T.Semester = 'S2000'
```

44

# Grouping

- But how do we compute the number of courses taught in S2000 *per professor*?

- Strategy 1: Fire off a separate query for each professor:

```
SELECT COUNT(T.CrsCode)
FROM Teaching T
WHERE T.Semester = 'S2000' AND T.ProfId = 123456789
```

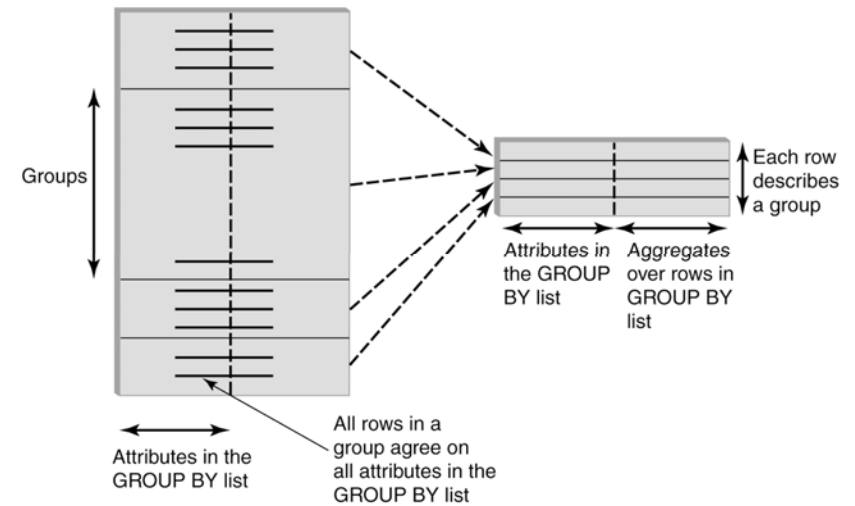
- Cumbersome
- What if the number of professors changes? Add another query?

- Strategy 2: define a special *grouping operator*:

```
SELECT T.ProfId, COUNT(T.CrsCode)
FROM Teaching T
WHERE T.Semester = 'S2000'
GROUP BY T.ProfId
```

45

# GROUP BY



46

# GROUP BY - Example

Transcript

1234	
1234	
1234	
1234	

Attributes:

- student's *Id*
- avg grade
- number of courses

```
SELECT T.StudId, AVG(T.Grade), COUNT (*)
FROM Transcript T
GROUP BY T.StudId
```

47

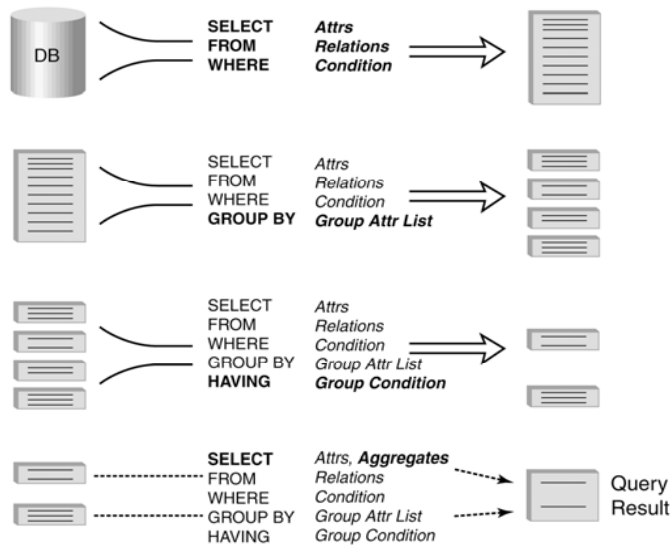
# HAVING Clause

- Eliminates unwanted groups (analogous to WHERE clause, but works on groups instead of individual tuples)
- HAVING condition is constructed from attributes of GROUP BY list and aggregates on attributes not in that list

```
SELECT T.StudId,
       AVG(T.Grade) AS CumGpa,
       COUNT (*) AS NumCrs
FROM Transcript T
WHERE T.CrsCode LIKE 'CS%'
GROUP BY T.StudId
HAVING AVG (T.Grade) > 3.5
```

48

## Evaluation of GroupBy with Having



49

## Example

- Output the name and address of all seniors on the Dean's List

```
SELECT S.Id, S.Name
FROM Student S, Transcript T
WHERE S.Id = T.StudId AND S.Status = 'senior'
```

```
GROUP BY < S.Id -- wrong
          S.Id, S.Name -- right
```

Every attribute that occurs in SELECT clause must also occur in GROUP BY or it must be an aggregate. S.Name does not.

```
HAVING AVG (T.Grade) > 3.5 AND SUM (T.Credit) > 90
```

50

## Aggregates: Proper and Improper Usage

```
SELECT COUNT (T.CrsCode), T.ProfId
-- makes no sense (in the absence of
GROUP BY clause)
```

```
SELECT COUNT (*), AVG (T.Grade)
-- but this is OK
```

```
WHERE T.Grade > COUNT (SELECT ....)
-- aggregate cannot be applied to result
of SELECT statement
```

51

## ORDER BY Clause

- Causes rows to be output in a specified order

```
SELECT T.StudId, COUNT (*) AS NumCrs,
       AVG(T.Grade) AS CumGpa
FROM Transcript T
WHERE T.CrsCode LIKE 'CS%'
GROUP BY T.StudId
HAVING AVG (T.Grade) > 3.5
ORDER BY DESC CumGpa, ASC StudId
```

Descending

Ascending

52

## Query Evaluation with GROUP BY, HAVING, ORDER BY

- As before {
- 1 Evaluate FROM: produces Cartesian product, A, of tables in FROM list
  - 2 Evaluate WHERE: produces table, B, consisting of rows of A that satisfy WHERE condition
  - 3 Evaluate GROUP BY: partitions B into groups that agree on attribute values in GROUP BY list
  - 4 Evaluate HAVING: eliminates groups in B that do not satisfy HAVING condition
  - 5 Evaluate SELECT: produces table C containing a row for each group. Attributes in SELECT list limited to those in GROUP BY list and aggregates over group
  - 6 Evaluate ORDER BY: orders rows of C

53

## Nulls

- *Conditions*:  $x \text{ op } y$  (where  $\text{op}$  is  $<$ ,  $>$ ,  $<=$ ,  $=$ , etc.) has value *unknown* (U) when either x or y is null
  - WHERE  $T.\text{cost} > T.\text{price}$
- *Arithmetic expression*:  $x \text{ op } y$  (where  $\text{op}$  is  $+$ ,  $-$ ,  $*$ , etc.) has value NULL if x or y is NULL
  - WHERE  $(T.\text{price}/T.\text{cost}) > 2$
- *Aggregates*: COUNT counts NULLs like any other value; other aggregates ignore NULLs

```
SELECT COUNT (T.CrsCode), AVG (T.Grade)
FROM Transcript T
WHERE T.StudId = '1234'
```

54

## Nulls (cont'd)

- WHERE clause uses a *three-valued logic* – T, F, U(*undefined*) – to filter rows. Portion of truth table:

C1	C2	C1 AND C2	C1 OR C2
T	U	U	T
F	U	F	U
U	U	U	U

- Rows are discarded if WHERE condition is F(*alse*) or U(*nknown*)
- Ex: WHERE  $T.\text{CrsCode} = \text{'CS305'}$  AND  $T.\text{Grade} > 2.5$

55

## Modifying Tables – Insert

- Inserting a single row into a table
  - Attribute list can be omitted if it is the same as in CREATE TABLE (but do not omit it)
  - NULL and DEFAULT values can be specified

```
INSERT INTO Transcript(StudId, CrsCode, Semester, Grade)
VALUES (12345, 'CSE305', 'S2000', NULL)
```

56

## Bulk Insertion

- Insert the rows output by a SELECT

```
CREATE TABLE DeansList (  
    StudId      INTEGER,  
    Credits     INTEGER,  
    CumGpa      FLOAT,  
    PRIMARY KEY StudId )  
  
INSERT INTO DeansList (StudId, Credits, CumGpa)  
SELECT      T.StudId, 3 * COUNT (*), AVG(T.Grade)  
FROM        Transcript T  
GROUP BY    T.StudId  
HAVING      AVG (T.Grade) > 3.5 AND COUNT(*) > 30
```

57

## Modifying Tables – Delete

- Similar to SELECT except:
  - No project list in DELETE clause
  - No Cartesian product in FROM clause (only 1 table name)
  - Rows satisfying WHERE clause (general form, including subqueries, allowed) are deleted instead of output

```
DELETE FROM Transcript T  
WHERE T.Grade IS NULL AND T.Semester <> 'S2000'
```

58

## Modifying Data - Update

```
UPDATE Employee E  
SET      E.Salary = E.Salary * 1.05  
WHERE    E.Department = 'R&D'
```

- Updates rows in a single table
- All rows satisfying WHERE clause (general form, including subqueries, allowed) are updated

59

## Updating Views

- Question: Since views look like tables to users, can they be updated?
- Answer: Yes – a view update changes the underlying base table to produce the requested change to the view

```
CREATE VIEW CsReg (StudId, CrsCode, Semester) AS  
SELECT      T.StudId, T. CrsCode, T.Semester  
FROM        Transcript T  
WHERE        T.CrsCode LIKE 'CS%' AND T.Semester='S2000'
```

60

## Updating Views - Problem 1

```
INSERT INTO CsReg (StudId, CrsCode, Semester)  
VALUES (1111, 'CSE305', 'S2000')
```

- **Question:** What value should be placed in attributes of underlying table that have been projected out (e.g., *Grade*)?
- **Answer:** NULL (assuming null allowed in the missing attribute) or DEFAULT

61

## Updating Views - Problem 2

```
INSERT INTO CsReg (StudId, CrsCode, Semester)  
VALUES (1111, 'ECO105', 'S2000')
```

- **Problem:** New tuple not in view
- **Solution:** Allow insertion (assuming the WITH CHECK OPTION clause has not been appended to the CREATE VIEW statement)

62

## Updating Views - Problem 3

- Update to a view might *not uniquely* specify the change to the base table(s) that results in the desired modification of the view (ambiguity)

```
CREATE VIEW ProfDept (PrName, DeName) AS  
SELECT  P.Name, D.Name  
FROM    Professor P, Department D  
WHERE   P.DeptId = D.DeptId
```

63

## Updating Views - Problem 3 (cont'd)

- Tuple <Smith, CS> can be deleted from ProfDept by:
  - Deleting row for Smith from Professor (but this is inappropriate if he is still at the University)
  - Deleting row for CS from Department (not what is intended)
  - Updating row for Smith in Professor by setting *DeptId* to null (seems like a good idea, but how would the computer know?)

64



## Updating Views - Restrictions

- Updatable views are restricted to those in which
  - No Cartesian product in FROM clause
  - no aggregates, GROUP BY, HAVING
  - ...

For example, if we allowed:

```
CREATE VIEW AvgSalary (DeptId, Avg_Sal) AS
  SELECT  E.DeptId, AVG(E.Salary)
  FROM    Employee E
  GROUP BY E.DeptId
```

then how do we handle:

```
UPDATE AvgSalary
  SET Avg_Sal = 1.1 * Avg_Sal
```

65

## Triggers and Active Databases

66

## Trigger Overview

- Element of the database schema
- General form:
  - ON *<event>* IF *<condition>* THEN *<action>*
  - *Event*- request to execute database operation
  - *Condition* - predicate evaluated on database state
  - *Action* – execution of procedure that might involve database updates
- Example:
  - ON updating maximum course enrollment
  - IF number registered > new max enrollment limit
  - THEN deregister students using LIFO policy

67

## Trigger Details

- **Activation** - Occurrence of the *event*
- **Consideration** - The point, after activation, when *condition* is evaluated
  - Immediate or deferred (when the transaction requests to commit)
  - *Condition* might refer to both the state before and the state after *event* occurs

68

## Trigger Details

- **Execution** – point at which *action* occurs
  - With deferred consideration, execution is also deferred
  - With immediate consideration, execution can occur immediately after consideration or it can be deferred
    - If execution is immediate, execution can occur before, after, or instead of triggering event.
    - Before triggers adapt naturally to maintaining integrity constraints: violation results in rejection of event.

69

## Trigger Details

- **Granularity**
  - *Row-level granularity*: change of a single row is an event (a single UPDATE statement might result in multiple events)
  - *Statement-level granularity*: events are statements (a single UPDATE statement that changes multiple rows is a single event).

70

## Trigger Details

- **Multiple Triggers**
  - How should multiple triggers activated by a single event be handled?
    - Evaluate one condition at a time and if true immediately execute action or
    - Evaluate all conditions, then execute actions
  - The execution of an action can affect the truth of a subsequently evaluated condition so the choice is significant.

71

## Triggers in SQL:1999

- **Events**: INSERT, DELETE, or UPDATE statements or changes to individual rows caused by these statements
- **Condition**: Anything that is allowed in a WHERE clause
- **Action**: An individual SQL statement or a program written in the language of Procedural Stored Modules (PSM) (which can contain embedded SQL statements)

72

## Triggers in SQL:1999

- **Consideration:** *Immediate*
  - Condition can refer to both the state of the affected row or table before *and* after the event occurs
- **Execution:** *Immediate* – can be before or after the execution of the triggering event
  - Action of before trigger cannot modify the database
- **Granularity:** Both *row-level* and *statement-level*

73

## Before Trigger Example

(row granularity)

Check that  
enrollment  $\leq$  limit

```
CREATE TRIGGER Max_EnrollCheck
  BEFORE INSERT ON Transcript
  REFERENCING NEW AS N --row to be added
  FOR EACH ROW
  WHEN
    ((SELECT COUNT (T.StudId) FROM Transcript T
      WHERE T.CrsCode = N.CrsCode
        AND T.Semester = N.Semester)
    >=
    (SELECT C.MaxEnroll FROM Course C
      WHERE C.CrsCode = N.CrsCode ))
  ABORT TRANSACTION
```

74

## After Trigger Example

(row granularity)

No salary raises  
greater than 5%

```
CREATE TRIGGER LimitSalaryRaise
  AFTER UPDATE OF Salary ON Employee
  REFERENCING OLD AS O
    NEW AS N
  FOR EACH ROW
  WHEN (N.Salary - O.Salary > 0.05 * O.Salary)
    UPDATE Employee -- action
    SET Salary = 1.05 * O.Salary
    WHERE Id = O.Id
```

Note: The action itself is a triggering event (but in this case a chain reaction is not possible)

75

## After Trigger Example

(statement granularity)

Keep track of salary  
averages in the log

```
CREATE TRIGGER RecordNewAverage
  AFTER UPDATE OF Salary ON Employee
  FOR EACH STATEMENT
  INSERT INTO Log
  VALUES (CURRENT_DATE,
    SELECT AVG (Salary)
    FROM Employee)
```

76